

facebook



# Thinking Outside the GIL

With AsyncIO and Multiprocessing

**John Reese**

Production Engineer, Facebook

@n7cmdr

[github.com/jreese](https://github.com/jreese)



# What's the GIL?

- Global Interpreter Lock
- One VM thread at a time
- No concurrent memory access
- I/O wait releases lock

# Stateful monitoring

- Gather ~100M data points
- Process and aggregate anomalies
- Easy to add new checks
- Simple deployment
- Few dependencies

```
def fetch(url):  
    return requests.get(url)
```

```
def process(url, response):  
    if 200 ≤ response.status ≤ 299:  
        ... # return anomaly
```

```
# store anomalies in sql somewhere
```



```
def fetch(url):  
    return requests.get(url)  
  
def process(url, response):  
    if 200 ≤ response.status ≤ 299:  
        ... # return anomaly  
  
# store anomalies in sql somewhere
```



# #impact

- One binary
- Fetch the world
- Process everything
- Aggregate results
- Thread pool for I/O



# Not aging well

- Scales in time and memory
- Runtime now too slow
- Underutilizing hardware
- Ultimately limited by the GIL

**Give me options**

# Switch to py3

~45% memory savings



~20% runtime reduction



# Sharding

- Technically correct
- Scales with number of workers
- Complicated deployments
- Communication overhead

# Multiprocessing

- Scales with CPU cores
- Automatic IPC
- Pool.map is really useful

```
def fetch(url):  
    return requests.get(url)
```

```
def fetch_all(urls):  
    with multiprocessing.Pool() as pool:  
        results = pool.map(fetch, urls)
```



```
def fetch(url):  
    return requests.get(url)
```

```
def fetch_all(urls):  
    with multiprocessing.Pool() as pool:  
        results = pool.map(fetch, urls)
```

# Multiprocessing

- Scales with CPU cores
- Automatic IPC
- Pool.map is really useful
- One task per process
- Beware forking, pickling

# AsyncIO

- Based on futures
- Faster than threads
- Massive I/O concurrency

```
async def fetch_url(url):  
    return await aiohttp.request("GET", url)
```

```
async def fetch_two(url_a, url_b):  
    future_a = fetch_url(url_a)  
    future_b = fetch_url(url_b)  
    a, b = await asyncio.gather(future_a, future_b)  
    return a, b
```

```
async def fetch_url(url):  
    return await aiohttp.request("GET", url)
```

```
async def fetch_two(url_a, url_b):  
    future_a = fetch_url(url_a)  
    future_b = fetch_url(url_b)  
    a, b = await asyncio.gather(future_a, future_b)  
    return a, b
```

```
async def fetch_url(url):  
    return await aiohttp.request("GET", url)  
  
async def fetch_two(url_a, url_b):  
    future_a = fetch_url(url_a)  
    future_b = fetch_url(url_b)  
    a, b = await asyncio.gather(future_a, future_b)  
    return a, b
```



```
async def fetch_url(url):  
    return await aiohttp.request("GET", url)  
  
async def fetch_two(url_a, url_b):  
    future_a = fetch_url(url_a)  
    future_b = fetch_url(url_b)  
    a, b = await asyncio.gather(future_a, future_b)  
    return a, b
```

# AsyncIO

- Based on futures
- Faster than threads
- Massive I/O concurrency
- Processing still limited by GIL
- Beware timeouts and queue length

**Why not both?**

# Multiprocessing + AsyncIO

- Use multiprocessing primitives
- Event loop per process
- Queues for work/results
- Highly parallel workload
- Need to do some plumbing

```
async def run_loop(tx, rx):  
    ...    # real work here
```

```
def bootstrap(tx, rx):  
    loop = asyncio.new_event_loop()  
    asyncio.set_event_loop(loop)  
    loop.run_until_complete(run_loop(tx, rx))
```

```
def main():  
    p = multiprocessing.Process(  
        target=bootstrap,  
        args=(tx, rx)  
    )  
    p.start()
```

```
async def run_loop(tx, rx):  
    ...    # real work here  
  
def bootstrap(tx, rx):  
    loop = asyncio.new_event_loop()  
    asyncio.set_event_loop(loop)  
    loop.run_until_complete(run_loop(tx, rx))  
  
def main():  
    p = multiprocessing.Process(  
        target=bootstrap,  
        args=(tx, rx)  
    )  
    p.start()
```



```
async def run_loop(tx, rx):  
    ...    # real work here  
  
def bootstrap(tx, rx):  
    loop = asyncio.new_event_loop()  
    asyncio.set_event_loop(loop)  
    loop.run_until_complete(run_loop(tx, rx))  
  
def main():  
    p = multiprocessing.Process(  
        target=bootstrap,  
        args=(tx, rx)  
    )  
    p.start()
```

```
async def run_loop(tx, rx):  
    limit = 10  
    pending = set()  
    while True:  
        while len(pending) < limit:  
            task = tx.get_nowait()  
            fn, args, kwargs = task  
            pending.add(fn(*args, **kwargs))  
  
        done, pending = await asyncio.wait(pending, ...)  
        for future in done:  
            rx.put_nowait(await future)
```



```
async def run_loop(tx, rx):
    limit = 10
    pending = set()
    while True:
        while len(pending) < limit:
            task = tx.get_nowait()
            fn, args, kwargs = task
            pending.add(fn(*args, **kwargs))

        done, pending = await asyncio.wait(pending, ...)
        for future in done:
            rx.put_nowait(await future)
```

```
async def run_loop(tx, rx):
    limit = 10
    pending = set()
    while True:
        while len(pending) < limit:
            task = tx.get_nowait()
            fn, args, kwargs = task
            pending.add(fn(*args, **kwargs))

        done, pending = await asyncio.wait(pending, ...)
        for future in done:
            rx.put_nowait(await future)
```

```
async def run_loop(tx, rx):  
    limit = 10  
    pending = set()  
    while True:  
        while len(pending) < limit:  
            task = tx.get_nowait()  
            fn, args, kwargs = task  
            pending.add(fn(*args, **kwargs))  
  
        done, pending = await asyncio.wait(pending, ...)   
        for future in done:  
            rx.put_nowait(await future)
```



```
async def run_loop(tx, rx):
    limit = 10
    pending = set()
    while True:
        while len(pending) < limit:
            task = tx.get_nowait()
            fn, args, kwargs = task
            pending.add(fn(*args, **kwargs))

        done, pending = await asyncio.wait(pending, ...)
        for future in done:
            rx.put_nowait(await future)
```

```
async def fetch_url(url):  
    return await aiohttp.request('GET', url)
```

```
def fetch_all(urls):  
    tx, rx = Queue(), Queue()  
    Process(  
        target=bootstrap,  
        args=(tx, rx),  
    ).start()  
  
    for url in urls:  
        task = fetch_url, (url,), {}  
        tx.put_nowait(task)  
  
    ...    # consume response queue
```

```
async def fetch_url(url):
    return await aiohttp.request('GET', url)

def fetch_all(urls):
    tx, rx = Queue(), Queue()
    Process(
        target=bootstrap,
        args=(tx, rx),
    ).start()

    for url in urls:
        task = fetch_url, (url,), {}
        tx.put_nowait(task)

    ...    # consume response queue
```

```
async def fetch_url(url):  
    return await aiohttp.request('GET', url)
```

```
def fetch_all(urls):  
    tx, rx = Queue(), Queue()  
    Process(  
        target=bootstrap,  
        args=(tx, rx),  
    ).start()  
  
    for url in urls:  
        task = fetch_url, (url,), {}  
        tx.put_nowait(task)  
  
    ...    # consume response queue
```



```
class Pool:
    async def queue(self, fn, *args, **kwargs) → int: ...
    async def result(self, id) → Any: ...
```

```
    async def map(self, fn, items):
        task_ids = [
            await self.queue(fn, (item,), {})
            for item in items
        ]

        return [
            await self.result(task_id)
            for task_id in task_ids
        ]
```



```
class Pool:
    async def queue(self, fn, *args, **kwargs) → int: ...
    async def result(self, id) → Any: ...

    async def map(self, fn, items):
        task_ids = [
            await self.queue(fn, (item,), {})
            for item in items
        ]

        return [
            await self.result(task_id)
            for task_id in task_ids
        ]
```



```
class Pool:
    async def queue(self, fn, *args, **kwargs) → int: ...
    async def result(self, id) → Any: ...

    async def map(self, fn, items):
        task_ids = [
            await self.queue(fn, (item,), {})
            for item in items
        ]

        return [
            await self.result(task_id)
            for task_id in task_ids
        ]
```

```
async def fetch_url(url):  
    return await aiohttp.request('GET', url)  
  
async def fetch_all(urls):  
    async with Pool() as pool:  
        results = await pool.map(fetch_url, urls)
```

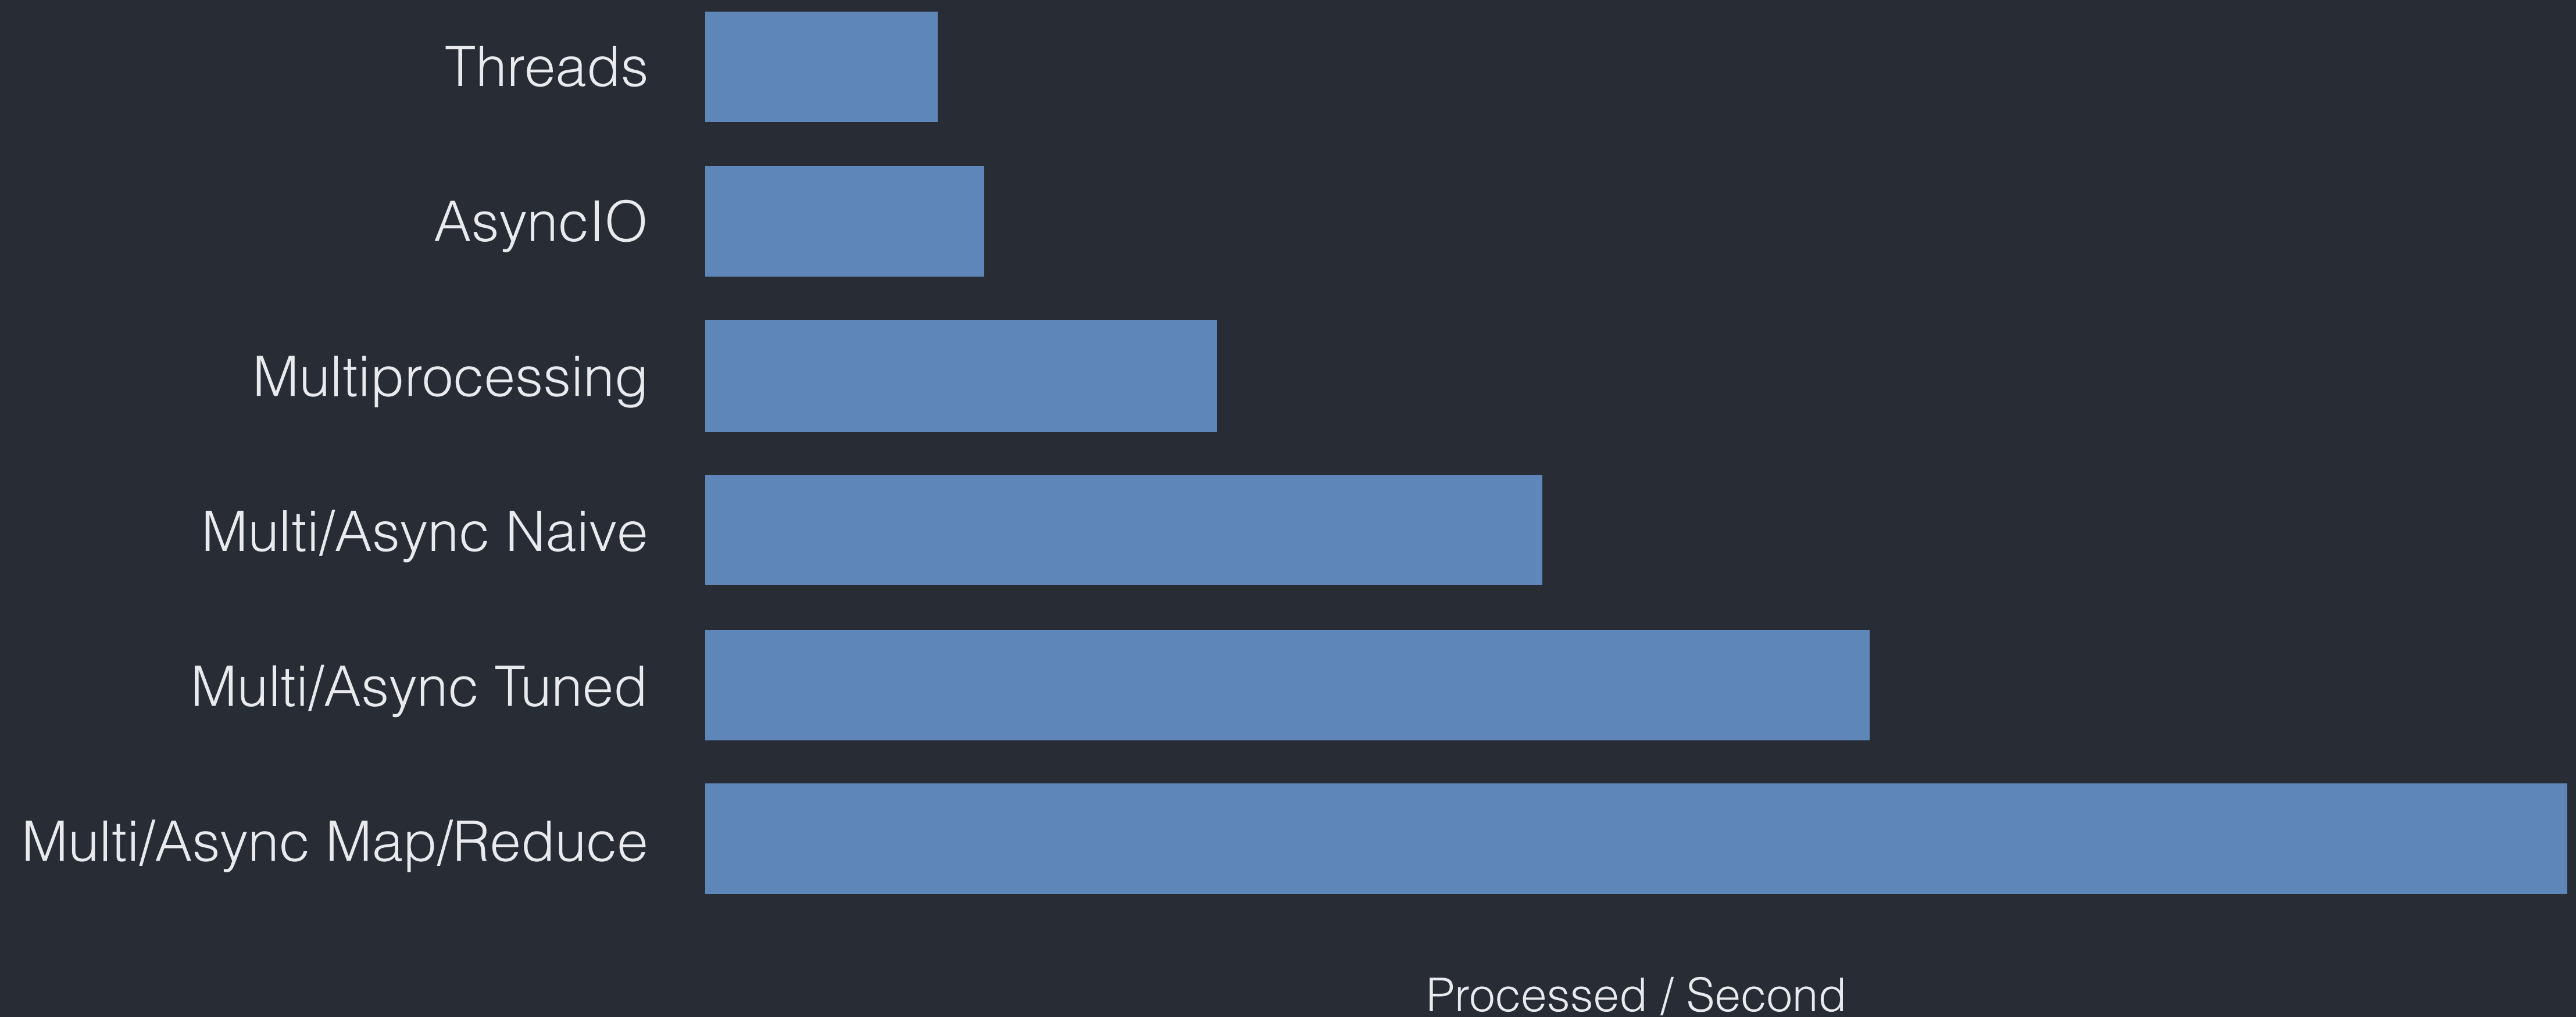
# Optimizations

- Multiple work queues
- Combine tasks into batches
- Use spawned processes

# Considerations

- Minimize what you pickle
- Prechunk work items
- Aggregate results in the child
- Use map/reduce

# Performance comparison



**I want it!**

```
$ pip install aiomultiprocess
```

# aiomultiprocess

[github.com/jreese/aiomultiprocess](https://github.com/jreese/aiomultiprocess)

- Simple implementation
- Emulates multiprocessing API
- One shot or process pool
- Supports map/reduce workloads



```
from aiomultiprocess import Pool

async def fetch_url(url):
    return await aiohttp.request('GET', url)

async def fetch_all(urls):
    async with Pool() as pool:
        results = await pool.map(fetch_url, urls)
```

**Python is *slow***

Python is ~~slow~~ powerful

**Great tools make  
complex tasks simple**

**John Reese**

Production Engineer, Facebook



@n7cmdr

[github.com/jreese](https://github.com/jreese)

facebook

